

Distributed and Parallel Computing

Lecture 08

Alan P. Sexton

University of Birmingham

Spring 2016

Floating Point Number Representation

Single Precision IEEE-754 floating point numbers has:

- 1 sign bit (S)
- an 8 bit biased exponent field with a bias of 127 (E)
 - 0 and 255 reserved for special numbers
 - $E=127$ corresponds to an exponent of $E - 127 = 0$
- a 23 bit mantissa field (M)
- for normal numbers, an implicit (hidden) initial 1 bit in the mantissa is assumed (24 binary bits \approx 7 decimal digits)
- a number of special bit patterns in S, E and M for positive and negative zeros, positive and negative infinities, NaN and sub-normal numbers

For normal numbers the interpretation is:

$$(-1)^S (1 + 2^{-23}M) 2^{E-127}$$

Subnormal Numbers

Normal numbers use the implicit 1 bit in the mantissa. What if only normal numbers are represented?

- The smallest representable number greater than 0 would be

$$(1 + 2^{-23} \times 0) 2^{-127} = 2^{-127} \approx 6 \times 10^{-39}$$

- The next smallest would be

$$(1 + 2^{-23} \times 1) 2^{-127} = 2^{-127} + 2^{-150}$$

- Thus much larger distance from 0 to $\text{succ}(0)$ than from $\text{succ}(0)$ to $\text{succ}(\text{succ}(0))$
- Sub-normal numbers: if exponent is -127 (i.e. $E=0$), don't use implicit 1 in mantissa

Rounding

- Unit in Last Place ($ULP(x)$): distance between the two floating point numbers that are the closest pair (a, b) that straddle x : $a \leq x \leq b$
- IEEE-754 requires that operations round to nearest representable floating point number: that the computed result be within 0.5 ULPs of the mathematically correct result.
- Consider $a + b$, where $a \gg b$, e.g. in a 4 digit decimal notation:

2.	3	4	5	
0.	0	0	1	2 3 4

Rounding

- Unit in Last Place ($ULP(x)$): distance between the two floating point numbers that are the closest pair (a, b) that straddle x : $a \leq x \leq b$
- IEEE-754 requires that operations round to nearest representable floating point number: that the computed result be within 0.5 ULPs of the mathematically correct result.
- Consider $a + b$, where $a \gg b$, e.g. in a 4 digit decimal notation:

2.	3	4	5				
0.	0	0	1	2	3	4	
<hr/>							
2.	3	4	6	2	3	4	

Sequence of Additions

Again in 4 digit decimal notation, consider summing a vector containing 1000.0 in the first element and then 10,000 elements of value 0.1.

- Mathematically, result should be 2000.0
- Incrementally adding from first returns 1000.0

$$\begin{array}{r} \boxed{1\ 0\ 0\ 0.}\ 0 \\ \quad\quad 0.\ \boxed{1\ 0\ 0\ 0} \\ \hline \boxed{1\ 0\ 0\ 0.}\ 0 \\ \quad\quad 0.\ \boxed{1\ 0\ 0\ 0} \\ \hline \vdots \end{array}$$

Worst case error in summing N numbers: $O(N)$

Fixing Sequence of Additions

To fix this, we can try adding in increasing order:

- Sort vector first
- Now incremental additions should add the smaller values together first then combine accumulated larger values with larger values from later in the vector

Problem:

- Accumulated small values may get significantly larger than later values in the vector
- Thus may help in some cases, but not a full solution

Kahan Sumation

Idea: Accumulate the sum, but calculate a correction term and add it to the next number at each step:


```
float fk_add(float * flt_arr)
{
    long i;
    float sum, correction, corrected_next_term, new_sum;

    sum = flt_arr[0];
    correction = 0.0;
    for (i = 1; i < ARR_SIZE; i++)
    {
        corrected_next_term = flt_arr[i] - correction;
        new_sum = sum + corrected_next_term;
        correction = (new_sum - sum) - corrected_next_term;
        sum = new_sum;
    }
    return sum;
}
```

- Worst case error in summing N numbers: $O(1)$, i.e. dependent only on the precision of the number representation
- But, not easy to parallelise

Fixing Sequence of Additions


Try sorting, then reducing



- adds similarly sized pairs in early steps
- but later steps add increasingly different sized values
- Worst case error in summing N numbers: $O(\log(N))$
- But: poor thread blocking, memory access patterns

Fixing Sequence of Additions

Can use more efficient compressed thread reduction pattern:



- but this adds values widely separated in vector
- Careful ordering helps
- Possible to use reduction for first few steps and switch to (unrolled) Kahan Summation to finish off

Absolute and Relative Errors

Given a value v , and its computed approximation \hat{v} :

- The *absolute error* in \hat{v} is $|v - \hat{v}|$
- The *relative error*, where $v \neq 0$, in \hat{v} is $\frac{|v - \hat{v}|}{|v|}$

The relative error is usually the more useful quantity.

- $|a| \gg |b| \Rightarrow a + b$ has a large absolute error
- $|b| \ll 1 \Rightarrow \frac{a}{b}$ has large relative and absolute errors
- $a \approx b \Rightarrow a - b$ has a large relative error (cancellation errors)

Example

We can compute $e = 2.7182818\dots$, the base of natural logarithms, with the formula:

$$e = \lim_{n \rightarrow \infty} \left(1 + \frac{1}{n}\right)^n$$

n	val
10^1	2.593742
10^2	2.704814
10^3	2.716924
10^4	2.718146
10^5	2.718268
10^6	2.718281
10^7	2.718282
10^8	2.718282
10^9	2.718282
10^{10}	2.718282
10^{11}	2.718282
10^{12}	2.718524
10^{13}	2.716110
10^{14}	2.716110
10^{15}	3.035035
10^{16}	1.000000
10^{17}	1.000000

Dealing with Roundoff errors

Rewrite formulae:

$$\frac{1}{\sqrt{x^2 + 1} - x}$$

Dealing with Roundoff errors

Rewrite formulae:

$$\frac{1}{\sqrt{x^2 + 1} - x}$$

When x is large, denominator has large error: possibly rounded to 0 resulting in divide by 0.

Dealing with Roundoff errors

Rewrite formulae:

$$\frac{1}{\sqrt{x^2 + 1} - x}$$

When x is large, denominator has large error: possibly rounded to 0 resulting in divide by 0.

$$\begin{aligned}\frac{1}{\sqrt{x^2 + 1} - x} &= \frac{\sqrt{x^2 + 1} + x}{(\sqrt{x^2 + 1} - x)(\sqrt{x^2 + 1} + x)} \\ &= \frac{\sqrt{x^2 + 1} + x}{x^2 + 1 - x^2} \\ &= \sqrt{x^2 + 1} + x\end{aligned}$$

Now tiny roundoff error

Dealing with Roundoff errors

Rewrite formulae:

$$\frac{1}{\cos^2 x - \sin^2 x}$$

Dealing with Roundoff errors

Rewrite formulae:

$$\frac{1}{\cos^2 x - \sin^2 x}$$

Normally fine when denominator is not near 0, but when x is near $\pi/4$, we get cancellation error

Dealing with Roundoff errors

Rewrite formulae:

$$\frac{1}{\cos^2 x - \sin^2 x}$$

Normally fine when denominator is not near 0, but when x is near $\pi/4$, we get cancellation error

$$\frac{1}{\cos^2 x - \sin^2 x} = \frac{1}{\cos 2x}$$

Now no problem where denominator is not near 0